

LEMON-Factor: Bidirectional Factor-Level Semantic Alignment for Graph–Text Pairs

Tomilov A.A.¹, Gineva D.², and Tirsikh D.²

¹ Independent researcher, affiliation to be confirmed

² ITMO University, St. Petersburg, Russia

Abstract. Graph–text evaluation has a common blind spot: a sentence may keep the right entity names while dropping the relation that makes a graph edge true. A faithful verbalization must preserve roles, direction, and predicate commitments; a reconstructed graph must recover them from text. We introduce LEMON-FACTOR, a bidirectional diagnostic framework that decomposes graph predicates into weighted semantic factors and evaluates both $\text{KG} \rightarrow \text{Text}$ coverage and $\text{Text} \rightarrow \text{KG}$ recoverability. The source graph remains the object of evaluation, but each predicate becomes an auditable set of roles, cues, polarity, direction, and evidence factors. On WebNLG, Forward LEMON-FACTOR reaches 0.7915 and Reverse LEMON-FACTOR reaches 0.7789. A MINE-STYLE comparison shows that node information is easier to recover than edge information. Controlled perturbations make the gap visible: deleting relation evidence reduces Forward LEMON-FACTOR by 0.211, while exact entity-label coverage changes by only 0.002. The same inventory format is applied to DrugProt and BC5CDR, where biomedical predicates make role, direction, polarity, and causality explicit. The study positions LEMON-FACTOR as an interpretable diagnostic layer for relation-level graph–text alignment; it is not a substitute for expert human validation.

Keywords: Graph-to-text evaluation · Semantic alignment · Knowledge graphs · WebNLG · Biomedical relation extraction · Predicate decomposition

1 Introduction

A graph edge is small, but it carries more than two names. In an edge (s, p, o) , the predicate p fixes the relation between the arguments: who did what, to whom, in which direction, and with which semantic commitment. If a graph-to-text system mentions the entities but loses the predicate, the sentence may still look fluent while no longer saying the same thing. The reverse problem is similar: a reconstructed graph is useful only if it recovers the relation, not merely the endpoints.

This failure is easy to overlook. A sentence can mention *Alan Bean* and *Wheeler, Texas* without stating that Wheeler is his place of birth. A biomedical abstract can preserve the same chemical and protein names while turning inhibition into activation. Standard WebNLG evaluation uses reference metrics such

as BLEU, METEOR, chrF++, TER, BERTScore, and BLEURT, together with human criteria such as data coverage, relevance, correctness, and fluency [1–3]. These signals remain useful, but they rarely identify which part of a predicate was preserved, weakened, or lost.

LEMON-FACTOR treats the predicate as the unit that must be audited. Each predicate is decomposed into typed, weighted semantic factors. For example, *birthPlace* contains a person-like subject, a biographical relation, and a place-like object. A DrugProt inhibition predicate contains a chemical actor, a gene/protein target, negative regulation, directionality, and evidence form. The score asks whether these factors are expressed in text or recovered in a reconstructed graph, and records which factors are missing.

The setting is deliberately restricted. We do not attempt general text–text similarity or open-ended semantic parsing. The source graph supplies the initial semantic structure, and LEMON-FACTOR audits the predicate layers that should remain stable across graph and text. In this sense, the method is a controlled graph–text instance of a broader layered-semantics idea: represent relation meaning explicitly, then compare it through interpretable parts rather than through a single opaque similarity score.

The framework is bidirectional. In $\text{KG} \rightarrow \text{Text}$, LEMON-FACTOR evaluates whether the text covers the semantic factors of the input graph. In $\text{Text} \rightarrow \text{KG}$, it evaluates whether a reconstructed graph preserves the source predicate factors. We also implement a MINE-STYLE node/edge baseline inspired by KGGen’s Measure of Information in Nodes and Edges (MINE), which evaluates text-to-KG information retention through recovered nodes and relations [19]. This baseline is a comparison point, not a reproduction of the full KGGen benchmark.

The experiments follow one line of reasoning: factor coverage beyond entity overlap on WebNLG, reuse of the same factors in the reverse direction, controlled perturbations for intended semantic damage, transfer of the inventory format to DrugProt and BC5CDR, and a small three-judge LLM probe for residual recoverability.

The contributions are a bidirectional factor-level formulation for graph–text semantic alignment; fixed factor inventories for WebNLG, DrugProt, and BC5CDR; reproducible forward, reverse, perturbation, and ablation analyses; a MINE-style node/edge comparison; and an exploratory LLM reliability probe. The paper is diagnostic rather than leaderboard-oriented: it shows when entity, triple, node/edge, vector-space, and factor-level views disagree, and what those disagreements reveal about relation-level semantic preservation.

2 Related Work

2.1 RDF-to-text benchmarks and reference metrics

The WebNLG Challenge established RDF-to-text generation as a benchmark task in which systems map DBpedia-derived RDF triples to natural language [1]. The WebNLG+ shared task extended the setting to bilingual and bidirectional generation and included both automatic and human evaluation [2].

The GEM benchmark argued for broader evaluation practice and challenge sets for NLG evaluation [3]. Recent LLM work on graph-to-text generation, including PlanGTG, has also renewed attention to planning, ordering, and attribution in graph-conditioned generation [4]. The standard WebNLG automatic stack includes reference-based metrics such as BLEU, METEOR, chrF++, TER, BERTScore, and BLEURT [5–7]. These metrics give useful sentence-level signals, but they do not localize an error to a source predicate or to a role inside that predicate.

2.2 Source-aware coverage and factuality metrics

A second line of work evaluates generated text against input data. PARENT weights n-gram precision and recall by evidence in the source table and reference text [8]. Data-QuestEval adapts question-answering evaluation to data-to-text semantic coverage without requiring a reference [9]. AlignScore learns a unified alignment function for factual consistency across several supervision sources [10]. The closest graph-to-text predecessor is FactSpotter, which decides whether a given triple (s, p, o) is present in a generated text and reports strong correlations with human judgments for data correctness, coverage, and relevance [11]. LEMON-FACTOR complements this line by moving from a binary whole-triple decision to a weighted decomposition of the predicate into semantic factors.

2.3 Atomic facts, semantic graphs, and directionality

Atomic factuality metrics decompose generated text into verifiable claims. FACTScore evaluates long-form text by decomposing it into atomic facts and checking support against a knowledge source [12]; VERIScore extends claim verification for long-form factuality [13]. This direction is text-driven: the output text is decomposed into claims. LEMON-FACTOR is source-driven: the source predicate is decomposed into factors that can be checked in text or in a reconstructed graph.

Semantic graph metrics such as Smatch compare graph structures, especially AMR graphs, by maximizing overlap under node alignment [14]. S²match and Smatch++ add softer or more standardized graph comparison variants [15, 16]. These metrics are central for parser evaluation, but they presuppose graph representations on both sides. LEMON-FACTOR instead evaluates factor evidence from graph-text pairs and only uses reconstructed graphs in the reverse experiment.

2.4 Text-to-KG extraction and MINE

Recent text-to-KG work has renewed interest in whether extracted graphs retain useful information from plain text. KGGen introduces a language-model-based text-to-KG generator and releases MINE, Measure of Information in Nodes and Edges, as a benchmark for testing whether an extractor produces useful KGs

from text [19]. Our MINE-STYLE baseline is inspired by this node/edge information split. It is not a reproduction of the KGGen benchmark, which targets KG extractors on plain text; rather, it adapts the idea to WebNLG graph-text pairs so that Text→KG recoverability can be compared with Reverse LEMON-FACTOR.

2.5 Behavioral testing and metric calibration

Evaluation metrics should be tested not only by correlation with human judgments, but also by whether they respond correctly to controlled changes. CheckList introduced behavioral tests for NLP models, including invariance tests, where label-preserving perturbations should not change predictions, and directional expectation tests, where perturbations should change predictions in a prescribed direction [20]. Perturbation CheckLists for NLG metrics apply the same idea to automatic metric evaluation [21]. He et al. use synthetic stress tests to expose blind spots of model-based NLG metrics [22]. Our calibration follows this methodology: relation deletion is a directional test, while entity aliasing and predicate paraphrase templates are invariance-style checks whose quality is explicitly marked.

Existing metrics therefore cover sentence-level similarity, source-aware coverage, binary triple checking, text-derived atomic facts, semantic graph overlap, and node/edge retention. LEMON-FACTOR targets the layer between a whole predicate label and a full graph: weighted factors inside source predicates, used in both KG→Text and Text→KG directions.

3 Method

3.1 Problem setting

Let a graph-text pair be (G, T) , where G is a set of directed edges $e = (s, p, o)$ and T is a verbalization. The forward task asks whether the text covers the information in the graph. The reverse task asks whether a graph reconstructed from the text, denoted $\hat{G}(T)$, preserves the information in the original graph. LEMON-FACTOR uses the same predicate-level semantic representation in both directions (Fig. 1).

3.2 Predicate factorization

Each predicate p is mapped to a weighted decomposition:

$$D(p) = \{(f_i, r_i, w_i)\}_{i=1}^n,$$

where f_i is a semantic factor, r_i is its role, and $w_i > 0$ is a salience weight. Roles include *predicate meaning*, *subject domain*, *object domain*, *value domain*, and modifier or background roles. The schema uses broad factors such as *person*, *place*, *organization*, *time*, *quantity*, and relation factors such as *biographical*




Fig. 1. Predicate-factor evaluation for graph-text alignment. A source predicate is decomposed into semantic factors, which are checked as forward graph-to-text coverage and reverse text-to-graph recoverability. Solid arrows denote forward evidence checks; dashed arrows denote reverse recovery checks.

relation, *location relation*, *part-whole relation*, *political relation*, and *creative relation*. This representation is intentionally coarser than a full ontology but richer than a single predicate label.

For example, *birthPlace* receives a person-like subject factor, a biographical predicate factor, and a place-like object factor. In the biomedical inventories, *chemical_inhibits_gene_or_protein* receives chemical and gene/protein role factors, an inhibition factor, negative polarity, directionality, and evidence-form factors. A text that mentions both entities but omits the relation cue should therefore receive partial credit, not full coverage. This is the key distinction from whole-triple binary scoring.

3.3 Forward LEMON: KG→Text coverage

For an edge e , text T , and component (f_i, r_i, w_i) , an evidence function $E(f_i, r_i, e, T) \in [0, 1]$ estimates whether the component is expressed. In the deterministic implementation, entity and domain factors are supported by normalized label or token evidence, while relation factors are supported by predicate cues and factor cues. The edge-level forward score is:

$$S_F(e, T) = \frac{\sum_i w_i E(f_i, r_i, e, T)}{\sum_i w_i}.$$

Example-level and corpus-level scores are averages over edges with available decompositions. The implementation stores covered and missing factors for each edge, so the scalar score is accompanied by a diagnostic trace.

3.4 Reverse LEMON: Text→KG recoverability

The reverse direction first constructs a deterministic reconstructed graph $\hat{G}(T)$ from the text using entity-label evidence and predicate lexical cues. Reverse LEMON-FACTOR then compares each gold edge with its best recovered counterpart. Subject and object recovery contribute through node/domain evidence;

predicate recovery contributes through overlap between the gold predicate’s factors and the recovered edge evidence. The reverse score therefore asks whether the source predicate’s semantic factors survive the text-to-graph reconstruction.

The current reconstructed graph is a lexical pseudo-extractor, not a full KG extraction model. This choice makes the experiment reproducible and keeps the focus on metric behavior rather than extractor quality. Future work can replace $\hat{G}(T)$ with KGGen, GraphRAG-style extraction, OpenIE, or domain-specific extractors.

3.5 MINE-style node/edge baseline

To compare factor-level recoverability with node/edge information retention, we implement a MINE-STYLE baseline inspired by KGGen’s MINE. Each gold triple is treated as a fact, and the reconstructed graph is searched for matching or related nodes and edges. The deterministic score separates node information, edge information, and a composite partial-credit node/edge score. We report it as a local baseline rather than a reproduction of the original benchmark.

3.6 Controlled perturbation calibration

A metric that measures semantic alignment should degrade under meaning-destroying perturbations and remain comparatively stable under meaning-preserving perturbations. We therefore create deterministic perturbations of both text and reconstructed graphs. Text-side destructive perturbations include relation phrase deletion, object swaps, and predicate swaps. Graph-side destructive perturbations include node drops, edge drops, and predicate swaps. Preserving perturbations include alias-like rewrites and template paraphrases, which are flagged as clean or weak-template depending on their reliability. All perturbed files and operation-level manifests are saved, enabling audit of each injected change.

4 Data and Experimental Resources

4.1 WebNLG sample

The WebNLG experiments use a stratified English sample built from RDF-to-text examples. The train split contains 200 examples and the development split contains 100 examples; both cover 16 WebNLG categories. The development set contains 308 graph edges. The sample is small by design: it keeps the complete scoring pipeline auditable while still exposing relation-level errors that are hidden by entity overlap.

4.2 Biomedical transfer resources

The biomedical transfer checks use two relation extraction corpora converted into the same graph-text JSONL format. DrugProt is a BioCreative VII resource for chemical-gene/protein relations with 13 relation types and 24,526

manually annotated relations [17]. In the current processed subset, the train split contains 500 abstracts, 3588 edges, 12 observed predicate types, and an average text length of 1729.5 characters. BC5CDR is a BioCreative V corpus of 1500 PubMed articles with chemical and disease annotations and 3116 chemical-disease interactions [18]. The converted BC5CDR train and development splits each contain 500 documents, with 985 and 961 CID edges respectively. These resources are used as transfer diagnostics only; we do not train or claim a new biomedical relation extraction system.

4.3 Predicate-factor inventories

The experiments use fixed JSON predicate-factor inventories. Each inventory contains a controlled factor vocabulary and one weighted decomposition per predicate. WebNLG uses 121 predicate decompositions and covers all 119 predicate labels observed in the evaluation files. DrugProt contains 13 biomedical relation decompositions and covers the 12 relation types observed in the processed subset. BC5CDR is intentionally a single-relation setting: its inventory decomposes the document-level chemical-induced disease relation into chemical exposure, disease outcome, causality, document-level evidence, and chemical-to-disease direction. This makes BC5CDR useful for testing whether the same factor format transfers to a biomedical relation with causality and document-level evidence, not for testing predicate-label diversity.

The inventories are deterministic resources. Earlier LLM-generated decomposition candidates are kept in the repository as development material, but they are not treated as expert gold and are not used as the main evidence in the deterministic scoring tables below. The validator checks that coverage is complete for the current evaluation files and that the weights of each predicate decomposition sum to one.

5 Experiments

The experiments test one representation from several angles. WebNLG provides an explicit RDF-to-text setting. DrugProt and BC5CDR then test whether the same factor format remains meaningful when predicates encode biomedical roles, direction, polarity, and causality. Deterministic scoring uses JSONL inputs and fixed JSON factor inventories. LLM calls appear only in the final reliability probe; they are not used to construct the deterministic scores.

5.1 Forward and reverse factor scoring

The first WebNLG experiment scores development texts using predicate decompositions and lexical cues. We run an error-driven resource expansion: missing predicates are identified, deterministic decompositions are added, lexical cues are expanded, and coverage is recomputed. This separates missing evaluation resources from missing information in the text.

To test bidirectionality, we reconstruct a graph from each WebNLG text with a deterministic lexical pseudo-extractor and compute Reverse LEMON-FACTOR. The reverse score is compared with expanded Forward LEMON-FACTOR. A small forward–reverse gap would indicate that the same factor schema can support both $\text{KG} \rightarrow \text{Text}$ coverage and $\text{Text} \rightarrow \text{KG}$ recoverability.

5.2 Node/edge retention and perturbations

We compare Reverse LEMON-FACTOR with a MINE-STYLE baseline inspired by KGGen’s Measure of Information in Nodes and Edges. The deterministic variant reports node information, edge information, and a composite node/edge score. A locked 50-fact subset shows how binary fact recoverability differs from a partial-credit node/edge composite. These values are diagnostic baselines, not a reproduction of the original KGGen/MINE benchmark.

Controlled perturbations are then injected into text and reconstructed graphs. The operators keep the source graph fixed and change only the text, which lets us ask whether a metric reacts to intended semantic damage rather than to a change in the reference graph. The final scoring uses five operators: node deletion, relation-cue deletion, argument swap, polarity flip, and relation blur. Polarity flips are most informative for DrugProt and BC5CDR, where relation labels explicitly encode activation, inhibition, regulation, or causality.

We score the shared perturbation files with deterministic baselines: entity recall, label match, a binary triple-style match, a MINE-STYLE node/edge harmonic score, a label-only LEMON variant, and the full deterministic LEMON-FACTOR proxy based on the fixed factor inventories and perturbation metadata. The same scoring layer is used for the compact ablation study. An offline character n-gram vector baseline is reported as a lexical/topical control, not as a dense embedding benchmark.

5.3 Exploratory LLM factor judging

The LLM reliability probe asks a different question: whether model judges can still recover the same factors after deterministic perturbations mark them as damaged. We sample perturbed instances from WebNLG, DrugProt, and BC5CDR and ask three independent LLM judges to label each semantic factor as covered, partially covered, or absent. The judges are Gemini 2.0 Flash, GPT-4o-mini, and Llama-3.1-70B. The prompts use JSON-schema constrained outputs through OpenRouter. The probe is reported as diagnostic evidence, not as human validation.

6 Results

6.1 From entity overlap to relation factors

The first WebNLG pass shows that predicate resources determine whether relation-level evaluation is possible. With the initial factor dictionary, only 140 of 308

development edges could be scored, and the missing decomposition rate was 0.5455. After deterministic decomposition and cue expansion, all 308 edges had predicate factors and Forward LEMON-FACTOR rose from 0.3683 to 0.7915. This increase is not a claim that all WebNLG texts are faithful. It shows that a factor-level metric needs explicit predicate resources before it can diagnose text quality.

Lexical controls isolate the added signal. Exact entity-label overlap is already high at 0.9215 and does not change under predicate-resource expansion. Token cosine and Jaccard are fixed because the graph-text strings do not change. The factor score changes because it checks relation roles, predicate meanings, and lexical evidence, not just entity mentions.

The reverse run gives a second check on the same representation. A deterministic lexical pseudo-extractor recovered 390 of 408 nodes and 196 of 308 edges from the WebNLG texts. Reverse LEMON-FACTOR reached 0.7789, close to the expanded Forward score of 0.7915. This does not evaluate a general-purpose extraction system; it shows that the same predicate factors can be used coherently in both directions.

The MINE-style comparison shows where the difficulty lies. Its deterministic node/edge composite is 0.8028, but the aggregate is dominated by node recovery. On the locked 50-fact subset, node information is 0.9600, while edge information and deterministic fact recoverability are both 0.5400. The pattern is consistent: entity mentions are easier to preserve and recover than the relations that connect them.

A relation-deletion sanity check makes the same point more directly. Removing relation evidence changes exact entity-label coverage by only -0.0020, while Forward LEMON-FACTOR drops by -0.2106. Predicate cue coverage and MINE fact recoverability fall by -0.3636 and -0.3539. Preserving the entities is therefore not the same as preserving the predicate; relation-aware scores are needed to see this failure mode.

6.2 Cross-domain perturbation sensitivity

The cross-domain perturbation files ask which metrics react when controlled edits damage relation meaning in WebNLG, DrugProt, and BC5CDR. Before reporting scalar drops, Table 1 summarizes the semantic layers that the current factor schema can audit. This layer view is intentionally limited: a perturbation can only be diagnosed when the corresponding factor is present in the inventory.

Figure 2 restates the same comparison as an annotated matrix. This view is easier to read than the earlier polar chart because the scale is linear and the cell values are explicit. It also makes the central point harder to miss: MINE-style and triple-style scores are harsher scalar detectors, while LEMON-FACTOR trades some scalar drop for a response that remains attached to explicit predicate factors.

Table 2 reports the same mean drops. Small values are not missing computations; they indicate weak sensitivity to that kind of damage. Large values should

Table 1. Layer view of the current factor schema. The rows are diagnostic layers, not independent models.

Layer	What the layer tests	Scope and diagnostic use
Entity/domain	Subject/object mention and coarse type support; stressed by node deletion.	Separates entity survival from relation preservation.
Predicate cue	Lexical or factor evidence for the relation; stressed by relation cue deletion and relation blur.	Detects when entities remain but the relation is no longer expressed.
Role/arg.	Which participant fills which predicate role; stressed by argument swap.	Identifies role reversals that entity-level scores may miss.
Direction	Directed effect or source-target orientation.	Applies when the inventory encodes directional predicates, mainly in biomedical relations.
Polarity	Activation, inhibition, causality, or negative regulation; stressed by polarity flip.	Applies only when polarity is explicit in the predicate inventory.
Evidence	Residual textual support for a source factor.	Separates intended perturbation damage from remaining contextual recoverability.

Table 2. Perturbation sensitivity as mean score drop. Larger values indicate that a metric reacts more strongly to the perturbation; unchanged or small drops mark blind spots rather than missing computations.

Perturbation	Entity	Triple	Node/edge	Label-only	LEMON-Factor
Node deletion	0.326	0.675	0.657	0.323	0.352
Edge deletion	0.215	0.655	0.637	0.482	0.281
Argument swap	0.214	0.540	0.521	0.321	0.413
Polarity flip	0.215	0.546	0.526	0.337	0.414
Relation blur	0.216	0.568	0.549	0.385	0.281


also be read carefully: they show a strong scalar reaction, not necessarily a more informative diagnosis.

The pattern is simple. Entity recall drops least because entity strings often remain when the relation is damaged. Triple-style and node/edge scores react more strongly, but collapse the error into one coarse decision. LEMON-FACTOR is most useful as a diagnostic decomposition: it reacts to argument swap and polarity flip when those dimensions are encoded in the inventory, and stays quiet when they are absent. The metric can diagnose only the predicate dimensions that the inventory makes explicit.

The ablation results identify what drives sensitivity. Removing participant-role factors produces the largest positive gain for the full model over the ablated variant, so roles are the main source of sensitivity to argument and node-level damage. Evidence, polarity, and direction have smaller effects, mostly in biomedical predicates where these factors are explicit. Label-only scoring can produce a larger scalar drop when surface labels disappear, but it cannot say whether the damage concerns roles, direction, polarity, or evidence.

6.3 LLM factor judging as a complementary view

The LLM probe evaluates semantic recoverability rather than deterministic damage. Each judge saw the source edge, factor inventory, and perturbed text, then



Higher values indicate stronger scalar sensitivity to induced damage; they do not by themselves indicate a better metric.

Fig. 2. Diagnostic perturbation matrix under controlled edits. Cell values are mean score drops, not absolute accuracy. Higher values indicate stronger scalar sensitivity to induced damage. MINE-style and triple-match act as coarse edge/fact detectors; vector cosine is the offline character n-gram vector baseline. LEMON-FACTOR is less harsh in aggregate, but its response remains tied to roles, polarity, direction, and evidence when those dimensions are encoded.

labeled each factor as covered, partial, or absent. Table 4 summarizes the three judge batches.

The mean LLM factor score is 0.809, pairwise agreement is 0.597, and agreement with deterministic perturbation labels is 0.561. This moderate agreement is useful rather than surprising. Deterministic perturbations record intended damage; LLM judges estimate how much of the original factor remains recoverable from residual context. LEMON-FACTOR exposes relation-level failure modes, while model-based judging can show which failures remain inferable.

7 Error Analysis

Two error classes recur across the experiments. The first is resource incompleteness. Predicates without decompositions, or with weak lexical cues, cannot be scored reliably. The error-driven expansion step exposes this problem before it contaminates the metric. After expansion, all 308 development edges had decompositions, but fallback decompositions remain less informative than expert-designed ones.

The second class is relation recoverability. Node information is high because WebNLG verbalizations often repeat entity labels, but edge information is substantially lower. The same pattern appears in the reverse experiment, in the MINE-style node/edge gap, and in the relation-deletion perturbation. The hard part of graph-text alignment is often not naming entities; it is preserving and detecting the relation that links them.

Table 3. Ablation sensitivity. Mean drop is the average perturbation response. Gain is full drop minus ablated drop; positive values mean that removing the component reduces sensitivity.

Variant	Mean drop Gain vs. full	
LEMON-full	0.320	0.000
Label-only	0.417	-0.097
Unweighted	0.310	0.010
Minus roles	0.185	0.135
Minus direction	0.310	0.010
Minus polarity	0.304	0.016
Minus evidence	0.290	0.030

Table 4. Exploratory LLM reliability probe. Judges are Gemini 2.0 Flash, GPT-4o-mini, and Llama-3.1-70B. Mean score treats covered, partial, and absent factor judgments as 1, 0.5, and 0; agreements are computed only on overlapping factor judgments.

Probe	Items	Judges	Mean score	Pairwise agr.	Det. agr.
LLM factor judging	45	3	0.809	0.597	0.561

The factor trace is the practical output of the method. A low LEMON-FACTOR can be traced to a predicate, a factor role, a missing cue, or weak subject/object evidence. The score is therefore not just a number at the end of an evaluation run. It is a debugging signal for datasets, generators, and extraction pipelines.

8 Limitations

The study does not yet include expert human validation. The factor inventories are deterministic resources built from predicate labels, schema constraints, rules, and development-time candidates. The LLM probe adds a semantic recoverability check, but model judgments are not a replacement for expert annotation. Prior LLM-as-judge work shows that model judges can scale evaluation while also introducing preference, calibration, and reasoning biases; the next version should add expert factor annotations and inter-annotator agreement [23, 24].

The deterministic evidence layer is limited. It can find labels and known cues, but it can miss implicit relations, long-range entailment, and paraphrases outside the cue dictionary. A stronger implementation should add embedding, entailment, or QA-based evidence while preserving the factor-level audit trail.

The reverse pipeline uses a lexical pseudo-extractor rather than a full text-to-KG system. This makes the experiment controlled, but the reverse scores should be read as metric-behavior diagnostics. Similarly, the MINE-style baseline adapts KGGen’s node/edge information idea to graph–text pairs; it is not a reproduction of the original MINE benchmark.

The scope is narrower than general semantic similarity. LEMON-FACTOR assumes that a source graph, relation annotation, or predicate inventory is available. It is therefore a graph-text fidelity metric and a restricted case of layered semantic comparison, not a general-purpose text-text similarity metric. The current paper also does not evaluate dense embedding baselines for arbitrary sentence pairs and does not claim to replace embedding-based retrieval or semantic textual similarity systems.

The present factor inventories cover relation-level roles, direction, polarity, causality, and evidence where those dimensions are encoded by the predicate resources. They do not model pragmatic phenomena such as irony, implicature, presupposition, stance, or rhetorical force. A future version could add such layers, but they are outside the current empirical scope.

Finally, the perturbations are deterministic and auditable, but still template-like. They support construct-validity checks for relation deletion, argument swap, and polarity flip, but they do not prove robustness to natural paraphrase. DrugProt and BC5CDR are used to test transfer of the factor format to biomedical roles, directionality, polarity, and causality; we do not claim state-of-the-art biomedical relation extraction performance.

9 Conclusion

LEMON-FACTOR evaluates graph-text alignment by opening relation predicates into semantic factors. This changes the unit of diagnosis. Instead of asking only whether a triple is present or whether two entity names overlap, the framework asks which roles, relation meanings, directions, polarity cues, and evidence factors survive from graph to text and back to graph.

The experiments show why this unit matters. On WebNLG, the same factor representation supports both forward coverage and reverse recoverability. The MINE-style comparison shows that node information is much easier to recover than edge information, and the relation-deletion sanity check shows that entity-label coverage can remain almost unchanged while relation coverage drops. Across WebNLG, DrugProt, and BC5CDR, the perturbation analysis shows that the score reacts to the semantic dimensions represented in the inventory: roles, direction, polarity, causality, and evidence.

The ablation and LLM reliability analyses sharpen this picture. Role factors contribute most to sensitivity under controlled damage. Label-only scoring remains less interpretable because it cannot say which part of the predicate was lost. LLM judges add a complementary signal: they often recover factors from residual context even when a deterministic perturbation marks them as damaged. Deterministic factor scoring measures intended, auditable semantic damage; model-based judging can reveal residual recoverability.

The result is a diagnostic layer between whole-triple factuality metrics and text-driven atomic-fact evaluation. LEMON-FACTOR does not replace existing metrics. It explains relation-level failures that coarser scores collapse into a single number.

Next, the provisional predicate inventories should be validated by domain experts, with accept, partial, missing-factor, and direction-error rates reported separately. The lexical evidence layer should then be paired with stronger paraphrase, dense embedding, entailment, or QA signals while keeping the factor trace intact. The reverse experiment should also be repeated with stronger text-to-KG extractors, including KGGen-style and domain-specific pipelines. More distant extensions include modality, quantification, presupposition, and pragmatic effects; these layers require separate annotation and should not be inferred from the present results.

References

1. Gardent, C., Shimorina, A., Narayan, S., Perez-Beltrachini, L.: The WebNLG Challenge: Generating Text from RDF Data. In: INLG, pp. 124–133 (2017). <https://doi.org/10.18653/v1/W17-3518>
2. Castro Ferreira, T., et al.: The 2020 Bilingual, Bi-Directional WebNLG+ Shared Task: Overview and Evaluation Results. In: Proceedings of the 3rd International Workshop on Natural Language Generation from the Semantic Web (2020)
3. Gehrmann, S., Adewumi, T., Aggarwal, K., et al.: The GEM Benchmark: Natural Language Generation, its Evaluation and Metrics. In: Proceedings of the 1st Workshop on Natural Language Generation, Evaluation, and Metrics, pp. 96–120 (2021). <https://doi.org/10.18653/v1/2021.gem-1.10>
4. He, J., Yang, Y., Long, W., Xiong, D., Gutierrez Basulto, V., Pan, J.Z.: Evaluating and Improving Graph to Text Generation with Large Language Models. In: NAACL-HLT, pp. 10219–10244 (2025). <https://doi.org/10.18653/v1/2025.naacl-long.513>
5. Papineni, K., Roukos, S., Ward, T., Zhu, W.-J.: BLEU: A Method for Automatic Evaluation of Machine Translation. In: ACL, pp. 311–318 (2002)
6. Zhang, T., Kishore, V., Wu, F., Weinberger, K.Q., Artzi, Y.: BERTScore: Evaluating Text Generation with BERT. In: ICLR (2020)
7. Sellam, T., Das, D., Parikh, A.P.: BLEURT: Learning Robust Metrics for Text Generation. In: ACL, pp. 7881–7892 (2020)
8. Dhingra, B., Faruqui, M., Parikh, A., Chang, M.-W., Das, D., Cohen, W.W.: Handling Divergent Reference Texts when Evaluating Table-to-Text Generation. In: ACL, pp. 4884–4895 (2019). <https://doi.org/10.18653/v1/P19-1483>
9. Rebuffel, C., Scialom, T., Soulier, L., Piwowarski, B., Lamprier, S., Staliano, J., Scoutheeten, G., Gallinari, P.: Data-QuestEval: A Referenceless Metric for Data-to-Text Semantic Evaluation. In: EMNLP, pp. 8029–8036 (2021). <https://doi.org/10.18653/v1/2021.emnlp-main.633>
10. Zha, Y., Yang, Y., Li, R., Hu, Z.: AlignScore: Evaluating Factual Consistency with a Unified Alignment Function. In: ACL, pp. 11328–11348 (2023). <https://doi.org/10.18653/v1/2023.acl-long.634>
11. Zhang, K., Balalau, O., Manolescu, I.: FactSpotter: Evaluating the Factual Faithfulness of Graph-to-Text Generation. In: Findings of EMNLP, pp. 10025–10042 (2023). <https://doi.org/10.18653/v1/2023.findings-emnlp.672>
12. Min, S., Krishna, K., Lyu, X., Lewis, M., Yih, W.-T., Koh, P.W., Iyyer, M., Zettlemoyer, L., Hajishirzi, H.: FActScore: Fine-Grained Atomic Evaluation of Factual Precision in Long Form Text Generation. In: EMNLP, pp. 12076–12100 (2023). <https://doi.org/10.18653/v1/2023.emnlp-main.741>

13. Sun, H., et al.: VERISCORE: Evaluating the Factuality of Verifiable Claims in Long-Form Text Generation. arXiv:2406.19276 (2024)
14. Cai, S., Knight, K.: Smatch: An Evaluation Metric for Semantic Feature Structures. In: ACL (2013)
15. Opitz, J., Daza, A., Frank, A.: AMR Similarity Metrics from Principles. Transactions of the Association for Computational Linguistics 9, 522–538 (2021). https://doi.org/10.1162/tacl_a_00329
16. Opitz, J.: Smatch++: Standardized and Extended Evaluation of Semantic Graphs. In: Findings of EACL, pp. 1595–1608 (2023). <https://doi.org/10.18653/v1/2023.findings-eacl.118>
17. Miranda-Escalada, A., Mehryary, F., Luoma, J., et al.: Overview of DrugProt BioCreative VII track: quality evaluation and large scale text mining of drug-gene/protein relations. Database 2023, baad080 (2023). <https://doi.org/10.1093/database/baad080>
18. Li, J., Sun, Y., Johnson, R.J., et al.: BioCreative V CDR task corpus: a resource for chemical disease relation extraction. Database 2016, baw068 (2016). <https://doi.org/10.1093/database/baw068>
19. Mo, B., Yu, K., Kazdan, J., Cabezas, J., Mpala, P., Yu, L., Cundy, C., Kanatsoulis, C., Koyejo, S.: KGGen: Extracting Knowledge Graphs from Plain Text with Language Models. arXiv:2502.09956 (2025)
20. Ribeiro, M.T., Wu, T., Guestrin, C., Singh, S.: Beyond Accuracy: Behavioral Testing of NLP Models with CheckList. In: ACL, pp. 4902–4912 (2020). <https://doi.org/10.18653/v1/2020.acl-main.442>
21. Sai, A.B., Dixit, T., Sheth, D.Y., Mohan, S., Khapra, M.M.: Perturbation Check-Lists for Evaluating NLG Evaluation Metrics. In: EMNLP, pp. 7282–7296 (2021). <https://doi.org/10.18653/v1/2021.emnlp-main.575>
22. He, T., Zhang, J., Wang, T., Kumar, S., Cho, K., Glass, J.: On the Blind Spots of Model-Based Evaluation Metrics for Text Generation. In: ACL, pp. 12067–12097 (2023). <https://doi.org/10.18653/v1/2023.acl-long.674>
23. Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R., Zhu, C.: G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment. In: EMNLP, pp. 2511–2522 (2023). <https://doi.org/10.18653/v1/2023.emnlp-main.153>
24. Zheng, L., Chiang, W.-L., Sheng, Y., et al.: Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. In: NeurIPS (2023)